

Nikunj Khitha

+91 9540234616

njkhitha2003@gmail.com

nikunj.codenex.dev

Nikunj Khitha

Nikunj2003

Professional Summary

Applied AI Engineer who takes agent systems from stakeholder discovery to production. Shipped Slack-accessible AI agents to **20+ business functions**, delivered **10+** production MCP servers into a shared enterprise tool registry, and built the OpenTelemetry/Langfuse evaluation platform gating every prompt, model, agent, and tool change on measured accuracy, latency, and cost. Python, TypeScript, Go, Java/Spring AI.

Technical Skills

Applied AI, Agents & Evaluation: AI Agents, Agentic AI, Multi-Agent Systems, Model Context Protocol (MCP), Tool Calling, Structured Output, Agent Memory, RAG, GraphRAG, Knowledge Graph RAG, Vector Search, Prompt Engineering, Prompt Caching, Model Routing, LLM Evaluation, LLM-as-a-Judge, Golden Datasets, Langfuse, Ragas, DeepEval, promptfoo, OpenTelemetry, LangChain, LangGraph, LangChain4j, Spring AI, CrewAI, Graphiti, Human-in-the-Loop Workflows, AWS Bedrock, OpenAI, Claude, Gemini, LLaMA

Languages & Backend: Python, TypeScript, Go, Java, SQL, Bash, FastAPI, Spring Boot, Spring Security, Node.js, Gin, SQLAlchemy, Next.js, React, REST APIs, Microservices, SSE, RBAC, JWT, OAuth2, Stripe

Data & Storage: PostgreSQL, pgvector, Neo4j, MongoDB, Elasticsearch, Redis, MSSQL, Supabase, MinIO, S3, Apache Superset

Platform & Observability: Docker, Kubernetes, AWS, Azure, CI/CD, Jenkins, GitHub Actions, Fabric8, Kafka, LocalStack, LiteLLM, n8n, Grafana, Prometheus, pytest, scikit-learn, Nginx, Resilience4j

Professional Experience

ArmorCode

Dec 2025 – Present

Applied AI Engineer (SDE 1)

Gurugram, India

- Won ArmorCode's **AI Ninja Award** as its first-ever and youngest recipient, presented on the company's global platform.
- Shipped Slack-accessible AI agents to **20+ business functions** by co-building Agentic Office OS, translating ambiguous stakeholder requirements into agents, skills, MCP integrations, and human-in-the-loop approval workflows, then owning rollout, debugging, and enablement.
- Gated every prompt, model, and MCP tool change on measured accuracy, latency, and cost with an OpenTelemetry/Langfuse platform scoring **9 AI surfaces** on golden datasets via deterministic checks, scikit-learn metrics, Ragas, and LLM-as-a-judge graders validated at **Cohen's kappa ≥ 0.7** , enforced as Jenkins CI gates.
- Secured multi-tenant agent tool use org-wide by delivering **10+ production MCP servers** into the shared enterprise registry, validating **20 of 20** authorization checks across **3 permission tiers** with OAuth2/RBAC and audit attribution.
- Cut recurring LLM spend as sole maintainer of a LiteLLM gateway, tracing **57%** of cost across **50,000+ requests** to **10+** automations, then migrating models and enabling Bedrock prompt caching while flagging a rewrite worth a further **$\sim 95\%$** .
- Grounded those agents with a tenant-isolated knowledge-graph RAG layer over **1M+ entities** in Neo4j and pgvector, scoring retrieval accuracy, ranking quality, and context precision across five query modes.
- Built a codebase-search MCP service in Go over **8 product repositories** that cut CS escalations, then restored it after **100+ queries** returned zero content by isolating a Go schema defect and a client timeout mismatch, flagging **14** exposed connectors.
- Kept AI-drafted documentation reviewable rather than blind-published by shipping Quill, an Electron/React GitHub PR review app pairing a WYSIWYG editor with an embedded agent terminal, per-branch Git worktrees, and permission-aware delivery.
- Replaced AWS QuickSight with a self-hosted business-data layer, consolidating **4 production AWS accounts** and **10+ SaaS sources** into PostgreSQL and Apache Superset via cross-account S3 replication, then exposing it to agents over MCP.

ArmorCode

Jan 2025 – Nov 2025

Software Development Intern (Applied AI)

Gurugram, India

- Created the initial Java framework for ArmorCode's platform agent, owning its short- and long-term memory layers, later shipping platform memory on **Graphiti** temporal knowledge graphs, and porting **2 of 6** sub-agents from LangChain4j to Spring AI.
- Owned backend integrations for **5+ security tools**, including Black Duck, Snyk, and Checkmarx, on an AppSec platform aggregating findings across **130+ connectors**.
- Automated new-integration scaffolding with AI-assisted code generation, template engines, and AST parsing, cutting per-integration boilerplate setup time by **30%**.
- Designed an OpenAI-compatible LLM proxy in Go serving Gemini CLI, Codex, and Claude Code, with multi-account load balancing, health-aware provider fallbacks, streaming, and Redis-backed caching.

Xansr Media (AIKO)

Jun 2024 – Dec 2024

Software Development Intern (GenAI Specialist)

Remote

- Built Fantasy GPT, a multi-step SQL-RAG system on a fine-tuned in-house model with LangGraph and DeepEval quality gates, resolving **98%** of complex sports queries against live match data.
- Consolidated multi-source sports data into MS SQL via Python ETL pipelines powering Fantasy GPT's SQL-RAG retrieval.
- Integrated Azure Speech SDK into AIKO for voice-driven sports media, live commentary in **20+ languages**, and on-the-fly highlight reels, for a product presented at **IBC 2024**.
- Optimized Node.js and FastAPI microservices to improve API performance **40%** and cut deployment time **42%** with Docker and GitHub Actions.

Projects

CodeNex: AI Builder: *Java, Spring Boot, Spring AI, React, TypeScript, SSE, Kubernetes, MinIO, Stripe*

[Live](#) | [GitHub](#)

- Launched a live AI code-generation SaaS that turns natural-language prompts into full React apps, using Spring AI, SSE streaming, and MinIO/NFS-backed workspace persistence.
- Provisioned per-user Kubernetes preview environments via Fabric8 and Ingress, with token quotas, RBAC, autoscaling, and Stripe subscription billing.

Serenify: *React, TypeScript, Supabase, pgvector, Gemini AI, Vercel*

[Live](#) | [GitHub](#)

- Released an open-source AI wellness product pairing Gemini-powered chat with mood tracking, journaling, guided sessions, crisis-help routing, and pgvector-backed personalization.
- Enforced per-user data isolation at the database with granular controls over what the assistant may read, so personalization improves across sessions without exposing sensitive history.

Education

B.Tech in Computer Science Engineering — The NorthCap University, Gurugram, India

Aug 2021 – Jun 2025